




MODEL CONTEXT PROTOCOL

# What Is MCP, and What Does It Mean for Modern Data Architectures?



Despite widespread excitement, investment, and industry noise, the advantages of AI remain elusive for today's enterprises.

A recent MIT report found that, despite growing enterprise investment in agentic AI, **95% of initiatives failed to reach production or deliver measurable ROI**. The core issue was not model quality, but brittle workflows and poor integration with company systems, leaving only 5% of pilots extracting real business value.

Deploying AI successfully is a challenge in itself. While **78% of companies have implemented** at least some form of AI tool or application, the number of companies **driving value from AI has stayed stubbornly low at around 26%**.

Leveraging standalone LLMs or off-the-shelf agents is one thing. Building the connective tissue that allows agentic systems to pull context from enterprise data and source systems in order to be effective—that's another challenge entirely.

That's where Model Context Protocol (MCP) can change the game. A recent innovation by the AI company, **Anthropic**, MCP is the key to connecting AI to the source systems and data it needs to be effective.

So what is MCP, and what does it mean for enterprise data architecture?

**This guide will walk you through the protocol, how it works, and why it matters for the modern data enterprise.** We'll look at key use cases for this integrative technology and explore the impact of MCP for enterprises looking to derive value from agentic systems at scale.

# Table of Contents

04

CHAPTER 1  
What is MCP?

---

06

CHAPTER 2  
Why MCP Matters:  
Giving Agents Better Data

---

07

CHAPTER 3  
How MCP Works

---

09

CHAPTER 4  
The Benefits of an  
MCP-Ready Architecture

11

CHAPTER 5  
MCP in Action:  
Real-World Use Cases

---

13

CHAPTER 6  
Is Your Organization  
Ready for MCP?

---

14

CHAPTER 7  
The Future is Real-Time  
and Agentic

## What is MCP?

From off-the-shelf LLMs like ChatGPT, AI has come a long way. We're now at the stage where companies can deploy AI agents: essentially autonomous, intelligent systems able to act on their own to complete tasks, enhance data streams, or detect fraud. To be effective, agentic AI needs a better way to get the data it needs. A way that is real-time, secure, and doesn't disrupt the business.

### The Integration Bottleneck

Until now, connecting agentic systems to source data has required engineers to build a patchwork of custom, hard-coded integrations for every new tool and data source.

This process is slow, expensive, and inefficient at scale. What the industry needs is a universal standard. A common language for AI and data systems to communicate. This is the problem that a standardized context protocol was designed to solve.

The integration bottleneck is the single biggest blocker to realizing the full potential of agentic AI

in the enterprise. To solve it, enterprises needed a standardized interface that would allow any AI agent to talk to any data source, securely and efficiently.

That interface is the **Model Context Protocol (MCP)**.

Introduced by the AI company Anthropic, MCP is an open standard that enables developers to build secure, two-way connections between their data sources and AI agents. The architecture is straightforward: data sources are exposed through MCP servers, and AI agents act as clients that can securely connect to them.

Think of it like the USB-C of AI. Before USB-C, every device had its own unique, proprietary charger. Famously, [Apple was forced via a legal mandate to abandon their lightning chargers](#) in favor of the universal standard.

Now, a single, standardized system can connect and power everything from your laptop to your phone (and even your iPhone). MCP aims to do the same for AI, replacing the tangled mess of custom integrations with one clean, universal protocol.




Image source: [Edwin Lisowski](#)

## How is MCP Different from an API?

Application Programming Interfaces (APIs) have long been the primary way for different software solutions to talk to each other. However, most APIs are custom-built for a specific purpose. Connecting a new AI agent to ten different data sources often means building and maintaining ten separate, point-to-point API integrations. This is laborious, inefficient, and costly for engineering teams.

MCP isn't a replacement for APIs, but rather a standardized layer that sits above them. It provides a single, consistent protocol for communication. Instead of building ten custom keys for ten different locks, MCP gives you one master keycard that can be securely programmed to open the right doors.

## Is MCP the Same as RAG?

Retrieval-Augmented Generation (RAG) is a popular technique used to improve the accuracy of large language models. With RAG, when a user asks a question, the system first "retrieves" relevant information from a knowledge base (like a vector database) and provides them to the AI as context along with the original question.

While RAG is powerful (and highly relevant to enterprise AI), MCP solves a different problem. RAG is about providing *static* context to an AI before it generates a response. MCP is about creating a *live, two-way connection* between an AI agent and operational data sources.

The two can complement each other. For instance, an MCP-powered data stream can provide the real-time information that makes a RAG system dramatically more accurate and relevant, but they are not the same thing.

## Why Standardization Matters

The shift from custom integrations to a universal standard is a game-changer for enterprise AI. By eliminating the need to manually code every connection, MCP dramatically accelerates the speed at which new agentic systems can be developed, tested, and deployed.


This creates a more interoperable and competitive ecosystem. As a result, enterprises can easily connect the best agents to the right data, without being locked into a single vendor's architecture.

In some instances, MCP can be deployed alongside an external memory layer to provide long-term context persistence. Unlike traditional stateless LLM queries, this allows a model to retain session-specific or user-specific context across multiple interactions.

This has direct implications for multi-turn conversations, where AI is given a form of memory. For customer support bots, agents with session memory can handle complex issues more fluidly without requiring users to re-explain themselves.

By reducing hallucination and maintaining continuity, MCP improves both performance and trust in long-form or long-term AI use cases.

Now that we understand what MCP is and why it matters, let's look at why it matters.



## Why MCP Matters: Giving Agents Better Data

AI is no longer just an isolated tool for answering questions or generating text, it is becoming a team of autonomous 'agents' that can execute complex, multi-step tasks. To act intelligently within your business, these agents need one critical ingredient: context.

Without access to fresh, relevant data, an AI agent lacks the direction it needs to be effective. It can reason, but it has no awareness of the current business landscape.

It doesn't know a customer's most recent purchase, the real-time status of a supply chain, or the latest update to a patient's medical record. Its actions are based on stale information, making its decisions generic at best. Dangerously wrong at worst.

### The Problem with Stale Data

Many enterprises have already invested heavily in modern data architectures like a data lakehouse. These platforms are powerful tools for business intelligence, reporting, and training machine learning (ML) models on vast amounts of historical data. To populate them, data is typically replicated from operational systems in batches, often on an hourly or daily schedule. For analytics and model training, this delay is perfectly acceptable.

However, for agentic AI that needs to understand the business environment right now, this latency is unacceptable. Your lakehouse is a reflection of the past, not the present. Using it to power a real-time agent limits the agent's ability to respond with meaningful, actionable data. This new pattern of agentic AI requires a new architectural approach: a real-time, operational replica of your data, designed specifically to give agents a live, up-to-the-second understanding of the enterprise environment.

### The Risk of Direct Access

So, why not just connect the AI directly to real-time, live production databases? This approach introduces two critical risks.

First, it can **overload your source systems**. Unlike traditional applications, agentic systems are non-deterministic; it is difficult to predict exactly when or how often they will query a data source to complete a task. This unpredictability makes it dangerous to connect agents directly to production databases, which are the lifeblood of your business. An unexpected spike in high-frequency queries from an army of AI agents could easily slow down or even crash the critical applications your customers and employees depend on. While you can attempt to put guardrails in place to limit their activity, these are just band-aids on a fundamental architectural problem.

Second, it creates an **additional governance burden**. To safely allow direct access, you would need to create and manage a complex new layer of granular, agent-specific permissions for every production database. Without this constant oversight, you risk exposing sensitive customer information, financial records, or other confidential data, creating a serious compliance risk.



## How MCP Works

We've explored how MCP acts as a universal translator between AI agents and data sources. But how does it work in practice?

Imagine your enterprise data is a giant, highly-sensitive library.

An AI agent is a researcher who needs specific information to complete a task, but you can't let the researcher roam the library without guardrails. Instead, the researcher works around a *specific request system* (MCP) that ensures that everyone uses the same forms, formats, and workflows to interact with the library.

### MCP consists of two layers:

- **Data layer:** Defines the JSON-RPC based protocol for client-server communication, including lifecycle management, and core primitives, such as tools, resources, prompts and notifications.
- **Transport layer:** Defines the communication mechanisms and channels that enable data exchange between clients and servers, including transport-specific connection establishment, message framing, and authorization.

Conceptually the data layer is the inner layer, while the transport layer is the outer layer.

Source: [modelcontextprotocol.io/docs/getting-started/intro](https://modelcontextprotocol.io/docs/getting-started/intro)

The researcher (**the AI agent**) submits a request for a specific piece of information. The front desk (the **MCP server**) receives the request, verifies the researcher's credentials and routes the request to the library (the **data source**) where the agent can retrieve only the precise information that was requested.

The researcher then gets the exact data they need, without ever getting access to the restricted stacks. The researcher is happy, the library is in good order, and the system stays secure.

## A Real-World Workflow

This client-server architecture allows for flexible and dynamic interactions. Instead of hard-coding custom integrations, MCP standardizes how these connections are made. Here is a step-by-step breakdown of the process:




### Step 1: The Handshake

The **AI agent** (the client) initiates a connection to the **MCP server**. In our analogy, this is the moment the researcher shows the librarian their credentials.

The server acts as a centralized and secure gateway to a catalog of available tools and data sources, such as a customer relationship management (CRM) system, an inventory database, or an internal knowledge base.

### MCP Integration Framework (Clients ↔ Servers ↔ AI Host)



- **The technical bit:** The connection is typically established over a secure WebSocket or HTTP connection. During this initial handshake, the client presents authentication credentials, such as an OAuth token, which the server validates to confirm the agent's identity and access permissions.



### Step 2: Tool Discovery

Once connected, the agent asks the server for a list of available tools and data it is authorized to access. The server responds with a standardized list that details each tool's capabilities.

This dynamic discovery means the agent doesn't need to be pre-configured for every possible tool; it can ask the server what's available and how to use it on demand.

- **The technical bit:** The server returns a structured JSON object that conforms to the MCP specification. This object contains a schema detailing the available **tools**, their **methods**, the required **parameters** for each method, and the expected **return** values. This allows the client to dynamically construct valid requests without prior knowledge of the tool's specific API.



### Step 3: The Request and Response

The agent then sends a request to perform a specific action, like retrieving the purchase history for a particular customer.

Because MCP uses a standardized communication protocol, the server understands the request, validates it, and routes the agent request to the right data. This structured exchange is what makes the entire process fast and scalable.

- **The technical bit:** The agent constructs a JSON-RPC request, which is a lightweight remote procedure call protocol. The request specifies the **tool\_id** and **method** it wishes to invoke, along with any necessary **parameters** in a JSON object. The server parses this request, executes the command against the backend data source, and sends back a JSON-RPC response containing either the requested data in the **result** field or a structured **error** object if the request fails.

This structured, secure, and repeatable process is what allows MCP to break the integration bottleneck, enabling enterprises to finally connect AI to real-time data at scale. Now let's explore the benefits of MCP, and why it's worth exploring for your enterprise.

## Model Context Protocol (MCP) – Complete Workflow

  
AI Agent  
(The Researcher)

  
Secure Gateway

  
MCP Server  
(The Librarian)

  
Data Sources  
The Library

### The Three-Step MCP Process



#### Step 1: The Handshake

AI Agent presents credentials and establishes secure connection



#### Step 2: Tool Discovery

Server provides list of available tools and authorized capabilities



#### Step 3: Request & Response

Agent requests specific data; server fetches and returns structured results





# The Benefits of an MCP-Ready Architecture

We've seen how MCP works as a secure and standardized go-between for AI agents and data. But what are the tangible advantages of putting this kind of architecture in place? Adopting a protocol like MCP, especially when powered by a real-time data foundation, delivers four foundational benefits that accelerate an enterprise's journey to AI maturity.


## 1. Faster Deployment and Iteration

The single biggest obstacle to innovation is the time it takes to get from idea to execution. The same is true of AI deployment. Too often the months-long process of building custom data integrations holds companies back from implementing agentic intelligence.

But because MCP replaces this manual, point-to-point coding with a standardized protocol, it dramatically reduces the time it takes to connect a new agent to a new data source. This means teams can move from idea to a functional proof-of-concept in days or weeks, not months.

This accelerated velocity lets enterprises adopt a more agile approach to AI development, where teams can rapidly test new ideas, measure their impact, and iterate quickly based on real-world results.

### *Faster Deployment, Faster Results*




## 2. Interoperability Across Systems

In a world where we're inundated with custom integrations, you're often locked into the vendors and tools that you've already spent significant engineering resources connecting. An MCP-based architecture breaks this lock-in.

Because it's an open standard, you can connect any MCP-compliant agent to any MCP-compliant data source. This creates a flexible, future-proof ecosystem with room to evolve and expand.

For instance, if a new, more powerful AI agent comes to market, you can swap it into your workflow without having to rebuild your entire data stack from scratch. This allows you to always use the best tool for the job, push for constant innovation and prevent costly dependencies on a single vendor.

### *MCP Makes it Easy to Swap Agents and Sources*



## 3. More Accurate and Relevant AI Output

AI systems are only as effective as the data they're fed. When that data is stale, the agent's responses can be generic, irrelevant, or simply wrong. This can lead to AI "hallucinations" that erode trust and deliver no business value. Worse, they can harm a business' reputation through poor outcomes.

By providing a live, continuous stream of accurate data, an MCP-ready architecture ensures that agents are always working with the most current and complete information.

This doesn't just improve the technical accuracy of the AI's output; it makes the output relevant to the immediate context of the task, project, or business problem at hand. That means better decisions, more reliable automation, and a higher degree of trust in the agentic systems you employ.




Image source: [meetergo](#)

#### 4. Secure and Governed Access to Data


As we discussed in Chapter 1, giving AI agents direct access to production databases is inherently risky. MCP provides a solution by creating a secure, auditable gateway. Every request from an agent is authenticated and authorized by the MCP server, ensuring the agent only accesses the specific data it is permitted to see.

When combined with a real-time replication platform, enterprises can create a safe, read-only copy of their production data for AI agents to use, completely isolating critical systems from AI workloads. This is essential for maintaining compliance in highly regulated industries like finance and healthcare.

**Learn more** about how Striim's MCP AgentLink makes MCP fast, scalable, and safe.

## Deploying Autonomous, Reactive AI

In the best-case scenario, advanced agentic systems are not triggered manually; instead, they are designed to react autonomously to business events as they happen. This is the core principle of an Event-Driven Architecture (EDA), and it represents the final piece of the puzzle for truly intelligent automation. In an EDA, every meaningful business action, a customer adding an item to a cart, an inventory level dropping below a threshold, a sensor detecting an anomaly—these are published as “events” to a real-time message queue.



MCP is the bridge that connects AI agents directly to these event streams. Instead of waiting for a user to ask a question, an agent can subscribe to a stream of relevant events. This transforms the agent from a passive, tool into a proactive, autonomous system. It can now act on opportunities the moment they arise or mitigate threats before they escalate, all without human intervention. The combination of a real-time EDA and an MCP-ready data layer is what enables the highest level of AI maturity: a truly responsive, event-driven, and intelligent enterprise.

Taken together, these key benefits: security, interoperability, speed, and accuracy are the reasons MCP architecture is so powerful. With it, enterprises are equipped not just to experiment successfully with agentic AI, but to deploy it for real-world impact.

## MCP in Action: Real-World Use Cases

We've walked through the technical advantages of MCP, and how it works, and its key benefits. But its true value can be found in the impact it has on the business.

MCP, when powered by a real-time data layer, unlocks a new class of context-rich agentic AI applications that can drive operational efficiency, enhance customer experiences, and create new revenue streams.



### Customer Service: Proactive, Personalized Support

- **The Challenge:** An enterprise customer service team was under constant pressure to communicate with multiple customers at once and resolve issues quickly. Yet, support agents, both human and AI-powered, didn't have the complete, up-to-the-minute view of a customer's account. As a result, customers were left frustrated by repetitive conversations and slow resolution times.
- **How MCP Solves It:** The company's IT team leverages MCP to let a specialized AI agent securely access a customer's entire account history in real time. The agent can safely gain access to a customer's latest purchase, recent support tickets, current shipping status, and more. It can also query multiple systems (CRM, order management, logistics) simultaneously to get a unified view of the customer's situation, and report back what it finds. With this context, the AI agent helps out the customer support reps with accurate answers, proactively identify potential issues (like a delayed shipment), and even process a return or offer a discount.
- **The Business Impact:** Customers are more satisfied, the team faces lower operational costs, and the overall support experience feels both personal and efficient, without extra effort.




### Retail Banking: Customer360 "Next Best Product"

- **The Challenge:** A retail bank wanted its relationship managers to instantly see a customer's full profile and receive AI-driven recommendations for the "next best product." But with customer, account, and transaction data spread across multiple core banking systems (Oracle and SQL Server), exposing production systems directly to an AI agent was not an option.
- **How MCP Solves It:** Using a platform like Striim, the bank continuously replicates data from Oracle and SQL Server into an Azure PostgreSQL MCP-ready zone. In-stream enrichment joins customer and transaction data in real time, while sensitive information remains masked. Fraud signals can also be detected in-stream with tools like Forseer. This creates a live, unified view of each customer that the MCP server exposes safely to the AI agent. From there, the agent can securely query consolidated profiles, compute propensity scores, and return personalized product recommendations.
- **The Business Impact:** Relationship managers gain instant access to a customer's 360-degree profile and AI-driven "next best product" suggestions. This enables faster, more relevant conversations that improve customer satisfaction and drive cross-sell revenue. At the same time, core banking systems remain protected, regulatory compliance is maintained, and fraud detection is strengthened.



## Healthcare: Up-To-The-Minute Patient Care

- **The Challenge:** A healthcare company found that providing continuous, high-quality care for patients involved constant monitoring and strict adherence to treatment plans. While clinicians relied on periodic patient check-ins, this often led to delays in identifying and addressing critical health events.
- **How MCP Solves It:** The company connects an AI agent to real-time data streams from patient wearables (like heart rate monitors or glucose sensors) and electronic health records (EHRs). Through MCP, the agent securely accesses this sensitive data, identifies anomalies that deviate from a patient's baseline, and cross-references them with the patient's medical history and prescribed treatment plan—all before they evolve into a major health issue.
- **The Business Impact:** The system provides instant alerts to the clinical care team when a patient's vitals are abnormal. This approach encourages early intervention, a critical factor in preventing the need for hospitalization while improving patient outcomes. It also automates routine monitoring, freeing up clinicians to focus on more complex cases and delivering higher-value care.



## Manufacturing: A Smarter Factory Floor

- **The Challenge:** An enterprise manufacturer can't afford downtime. Critical failures lead to a halt to production, missed deadlines, and costly delays. And without real-time data over their supply chain, they're in the dark. They have no way of knowing if they have the right parts in the right place at the right time.
- **How MCP Solves It:** Bringing on an AI agent specifically for predictive maintenance, the company can deploy an intelligent system as the factory's central nervous system. Using MCP, they connect the agent to live data from sensors on every machine, listening for the changes in vibration or temperature that signal impending breakdown. The agent also taps into real-time inventory and shipping data, allowing it to not only predict that a specific part will fail next week, but to automatically order a replacement that arrives in time for a prompt repair.
- **The Business Impact:** Instead of reacting to problems after they've already cost time and money, they can fix them before they happen. This predictive approach keeps the production line moving and makes the entire supply chain smarter and more resilient. The result is less downtime, lower costs, and a manufacturing operation that runs like a well-oiled machine.

The most effective use case for MCP is the one that best aligns with your business priorities. The point is not to try and do everything at once, but to select one or two high-impact areas: testing, iterating, and measuring results before moving on.

## Is Your Organization Ready for MCP?

Taking the next step to adopt MCP architecture requires a baseline readiness across your data, governance, and infrastructure. So, how do you know if your organization is prepared to seize the opportunity?

### Signs an MCP-Ready Architecture Will Deliver Value

If you recognize your organization in the following statements, it's a strong signal that you are ready to benefit from building an MCP-ready data architecture.

- **You have a clear AI mandate, but are still grappling with execution.** Your leadership team has prioritized AI, but your technical teams are hitting a wall when it comes to connecting models to the data they need to be effective. The integration bottleneck is slowing you down.
- **Your business needs data to make decisions and operate effectively.** You operate in an industry like finance, logistics, or e-commerce, where data is critical to successful business outcomes. Delayed data costs you opportunities and creates significant risks.
- **Data governance and security are non-negotiable.** You handle sensitive customer, patient, or financial data, and the risk of exposing production systems to AI workloads is unacceptable. You need a way to provide AI with data that is both fresh and compliant.
- **You are investing in a multi-vendor, best-of-breed AI strategy.** You want the flexibility to use the best AI models and tools from different providers, but you are concerned about the complexity and cost of creating a fragmented web of custom integrations.

### Warning Signs: Foundational Work May Still Be Required

Conversely, there are a few signs that your organization may have work to do. You may need to address foundational data challenges before jumping straight in with MCP.

- **Your data is heavily siloed and inconsistent.** If different departments have their own versions of the truth and there is no unified view of core entities like "customer" or "product," an AI agent will only amplify the confusion.

- **You lack a clear data governance framework.** Without defined ownership, access policies, and quality standards, it is difficult and risky to securely expose data to any new system, including an AI agent.
- **Your current infrastructure struggles to deliver real-time data.** If your architecture is primarily built for daily or hourly batch processing, you will need to modernize your data stack to provide the low-latency data streams that make agentic AI effective.

### Your Next Steps

If you've identified that your organization is ready to move forward, the path to an MCP-ready architecture doesn't have to be a multi-year overhaul. Taking it step by step, it's possible to realize the benefits of MCP within a short timeframe.

1. **Start with a single, high-impact use case.** Don't try to boil the ocean. Identify one specific business problem, like the ones we explored in Chapter 5, where real-time, context-aware AI can deliver measurable value.
2. **Identify the data you need.** Map out the critical data sources required for that use case. This might include a production database, a SaaS application like Salesforce, and a streaming data source like Kafka.
3. **Build the real-time data layer.** This is the critical step. To power your MCP server, you need a way to continuously stream data from your sources into a safe, scalable, and AI-ready destination. This is where a real-time data integration and streaming platform becomes essential.

By focusing on a single use case and building the right data foundation, you can demonstrate the value of an MCP-based approach quickly, get buy-in from business leaders and technical teams, and build momentum for broader adoption across the enterprise.

## The Future is Real-Time and Agentic

The emergence of MCP is a breakthrough technical innovation. More than that, it's a signal of where the industry is heading.

The future of enterprise AI will not be confined to LLM chatbots or standalone applications. It's about embedding AI in the enterprise: building a deeply integrated ecosystem where autonomous agents can securely access the rich, real-time data they need to deliver meaningful outcomes.

As MCP becomes the standardized bridge between AI and data, the competitive advantage will shift. It will no longer be enough to simply have or use AI in the enterprise; the advantage will go to the organizations that can provide their agents with the freshest, most reliable, and most context-rich data.

### Striim AgentLink: Your Governed MCP-Ready Data Layer

This is where Striim provides the essential bridge to the future. MCP solves the integration problem, but it does not solve the underlying data problem. For an MCP server to be effective, it must be powered by a data layer that is real-time, secure, and context aware.

Striim's real-time intelligence platform is designed to create this MCP-ready data layer. [MCP AgentLink](#) offers a **continuous, real-time, cleansed, and protected operational replica** of your production systems in a safe, compliant zone. This gives agents fresh, accurate data without ever exposing or impacting your core operational databases.

MCP AgentLink allows you to:

- **Accelerate agent adoption** by giving AI agents instant, governed access to operational data without exposing production systems or slowing performance.
- **Improve decision quality** by streaming data that is not just real-time, but cleansed, governed, and contextually enriched so agents act on the right information, every time.
- **Reduce compliance risk** by embedding controls like sensitive data masking, vector enrichment, and workflow validation directly into the data stream, making governance automatic.

The path to scaling agentic AI is no longer a mystery. It requires a standardized protocol like MCP and a real-time data foundation to power it. With Striim, enterprises can build that foundation today, moving from the ambition of AI to the reality of intelligent, autonomous operations.

### Take the Next Step Towards Agentic AI

Ready to see how a real-time, MCP-ready data layer can accelerate your AI initiatives?

[Try Striim for Free](#), or [Request a Demo](#) to learn more.





## About Striim

Striim powers real-time intelligence for enterprise AI with intelligent data integration and event-driven streaming that operationalizes AI at scale. The platform delivers high-quality, decision-ready data for real-time inference, contextual reasoning in RAG, and event-driven agentic automation—processing over 100 billion events daily with industry-leading CDC, in-stream transformations, and sub-second latency. Its fully managed SaaS unifies data across clouds, applications, and databases with enterprise-grade availability and elastic scalability. With MCP-ready AgentBridge, Striim provides continuous, cleansed replicas in safe, compliant zones so agents get fresh, accurate data without exposing production systems. Intelligence is embedded directly into the stream, with AI agents generating vector embeddings, detecting anomalies, governing sensitive data, and validating pipelines for audit readiness. To learn more, visit [www.striim.com](https://www.striim.com).

To Learn More About Striim and MCP AgentLink,  
Visit [go2.striim.com/mcp-agentlink](https://go2.striim.com/mcp-agentlink)